

# Interval Tree

An interval tree is a data structure that is used for holding a collection of intervals in a way that allows for a fast test of what intervals overlap a point. By extending a binary search tree, this can be done quickly.

## ADT

The public interface of this interval tree is:

- `add(start, end, name)` : Add the given interval to the tree
- `remove(name)` : Remove the given interval from the tree
- `clear()` : Remove all items from the tree
- `getEndpoints(name)` : Return the start and end of a given interval
- `testPoint(point)` : Return a set of all of the intervals that overlap with a given point
- `testRange(start, end)` : Return a set of all the intervals that overlap with a given interval


All intervals in this tree (including the arguments to `testRange` ) are endpoint inclusive.

Each node in the tree stores that intervals that it covers, as well as a boundary value if it is not a leaf. Each node also stores the names of the intervals have been inserted into it.

The tree itself stores all of endpoints for the intervals contained in it. This is used for the `getEndpoints` method, but more importantly it is used to significantly optimize the removal method.

## Background


For a collection of intervals on the number line, the corresponding elementary intervals are created by breaking it up wherever one interval start or ends. An example of this is shown below.



Each node in the interval tree represents a subinterval. Non-leaf nodes also contain a boundary value. The nodes are arranged as a binary search tree with the boundaries as the keys. In this tree, the leaves are special, since they have no boundary and are not considered in checking the binary search tree property.

Given a node that represents the interval  $[x, z]$  and has the boundary value  $y$ , then the left child of that node represents the interval  $[x, y]$  and the right child represents  $[y, z]$ .

When an interval is stored in a tree, it is placed in the highest node which represents that interval. If necessary, an interval can be broken down and stretched across multiple nodes.



For example, to insert an interval named `a` covering  $[5, 10]$  to the above tree, the name `a` is just added to the yellow leaf node.

If you wanted to insert the interval `b` covering  $[5, 15]$  into this tree, the interval needs to be split up into  $[5, 10]$  and  $[10, 15]$ . This means that `b` is added to both the yellow and green leaf nodes.

An interesting note about this way of implementing an interval tree is that it can hold intervals with endpoints at infinity. To store the interval `c` covering  $(-\infty, 10]$ , you could split the interval up into  $(-\infty, 5]$  and  $[5, 10]$ , but the better solution is to store the interval in the node labelled 5, since this node represents exactly the interval that is being added. This same logic applies to nodes without infinite endpoints. If a node has the same interval in both of its children, that interval can more efficiently be placed in the parent node.

## Implementation

The interval tree class is very simple. For the most part, it just delegates function calls to the root node.

```
1 class IntervalTree:
2     def __init__(self):
3         self._root = IntervalNode()
4         self._dict = {}
5
6     def add(self, start, end, name):
7         if not (start < end):
8             raise ValueError("The start of an interval must \
9                               be smaller than its end")
10
11         if name in self._dict:
12             raise ValueError("Interval with same name already in tree")
13         self._root = self._root.add(start, end, name)
14         self._dict[name] = (start, end)
15
16     def remove(self, name):
```

```

16         if not name in self._dict:
17             raise KeyError("Interval not in tree")
18         start, end = self._dict[name]
19         self._root = self._root.remove(name, start, end)
20         del self._dict[name]
21
22     def clear(self):
23         self._root = IntervalNode()
24
25     def getEndpoints(self, name):
26         return self._dict[name]
27
28     def testPoint(self, point):
29         return self._root.testPoint(point)
30
31     def testRange(self, start, end):
32         return self._root.testRange(start, end)
33
34     def __repr__(self):
35         return repr(self._root)

```

There are some basic methods that the node class has. The node stores its intervals in a set to avoid duplicates and to make node manipulations easier.

If the node has a boundary, that means it's not a leaf. Leaves have both of their children set to `None` while non-leaf nodes have other nodes for both children.

`self.min` and `self.max` store the boundaries of the subinterval that the particular node represents.

```

1 class IntervalNode:
2     def __init__(self, boundary = None, intervals = (), min = ninf, max = inf):
3         self.boundary = boundary
4         self.intervals = set(intervals)
5         self.left = IntervalNode(min = min, max = boundary) \
6             if boundary is not None else None
7         self.right = IntervalNode(min = boundary, max = max) \
8             if boundary is not None else None
9         self.min = min
10        self.max = max
11        self._height = 0

```

There is also a function to nicely print out the tree.

```

1     def _nicerepr(self, n = 1):
2         ret = (('B : ' + repr(self.boundary)) if not self._isleaf else "L")\
3             + ' : ' + repr(list(self.intervals))
4         if self.left is not None:
5             ret += '\n' + ' ' * n + '< ' + self.left._nicerepr(n + 1)
6         if self.right is not None:
7             ret += '\n' + ' ' * n + '> ' + self.right._nicerepr(n + 1)
8         return ret
9
10    def __repr__(self):
11        return self._nicerepr()

```

Then just a function to check if a node is a leaf.

```

1     @property
2     def _isleaf(self):
3         return self.boundary == None

```

## Balancing

The next block contains methods used for balancing the tree using the AVL algorithm as shown in the textbook. Leaf nodes don't count towards the height of the tree because not having a boundary means they can't be rotated.

```

1     def _updateheight(self):
2         if self._isleaf:
3             return 0
4         lheight = -1 if self.left._isleaf else self.left._height
5         rheight = -1 if self.right._isleaf else self.right._height
6         return 1 + max(lheight, rheight)
7
8     def rebalance(self):
9         bal = self._balance
10        if bal == -2:
11            if self.left._balance > 0:
12                self.left = self.left._rotateleft()
13            newroot = self._rotateright()
14        elif bal == 2:
15            if self.right._balance < 0:
16                self.right = self.right._rotateright()
17            newroot = self._rotateleft()
18        else:
19            return self
20
21        newroot.left._height = newroot.left._updateheight()
22        newroot.left._height = newroot.right._updateheight()
23        newroot._height = newroot._updateheight()
24
25        return newroot
26
27    @property
28    def _balance(self):


```

```

29     rheight = self.right._height if self.right else 0
30     lheight = self.left._height if self.left else 0
31     return rheight - lheight

```

Rotation is a bit more complicated in this tree than in a normal binary search tree. As shown below, after a rotation the old root and new root now represent different subintervals. This means that those two nodes need to be updated. The intervals stored in those nodes then need to be moved so they stay in the correct node. In addition to this, it is not possible to rotate a leaf node to be a parent, since it has no boundary and can not be ordered according to the binary search tree property.



For a right rotation as shown in the diagram above, there are three groups of interval movements that need to be done:

1. The intervals stored in the old root (10) are moved to the new root (5)
2. The intervals stored in the new root are moved into both the red and green leaves
3. The intervals in both the green and blue nodes are combined and moved up into the old root (10).

A left rotation works similarly. After this, the heights of the old and new roots also need to be updated.

The code for the rotations, then, is

```

1     def _rotateright(self):
2         if self.left._isleaf:
3             raise RotationError("Can't move a leaf node up")
4
5         # Make copy of old intervals before clearing
6         oldrootintervals = set(self.intervals)
7         newrootintervals = set(self.left.intervals)
8
9         self.intervals.clear()
10        self.left.intervals.clear()
11
12        # A standard rotation
13        newroot = self.left
14        self.left = newroot.right
15        newroot.right = self
16
17        # Adjust subintervals represented by the nodes

```

```

18     newroot.min = self.min
19     newroot.max = self.max
20     self.min = newroot.boundary
21
22     # And the interval transfers mentioned above
23     newroot.intervals = oldroot.intervals
24
25     newroot.left.intervals |= newroot.intervals
26     self.left.intervals |= newroot.intervals
27
28     # Join up any intervals that are in both children of oldroot
29     both = self.left.intervals & self.right.intervals
30     self.intervals |= both
31     self.left.intervals -= both
32     self.right.intervals -= both
33
34     self._height = self._updateheight()
35     newroot._height = newroot._updateheight()
36     return newroot
37
38 def _rotateleft(self):
39     if self.right._isleaf:
40         raise RotationError("Can't move a leaf node up")
41
42     oldrootintervals = set(self.intervals)
43     newrootintervals = set(self.right.intervals)
44
45     self.intervals.clear()
46     self.right.intervals.clear()
47
48     newroot = self.right
49     self.right = newroot.left
50     newroot.left = self
51
52     newroot.min = self.min
53     newroot.max = self.max
54     self.max = newroot.boundary
55
56     newroot.intervals = oldrootintervals
57
58     newroot.right.intervals |= newrootintervals
59     self.right.intervals |= newrootintervals
60
61     both = self.left.intervals & self.right.intervals
62     self.intervals |= both
63     self.left.intervals -= both
64     self.right.intervals -= both
65
66     self._height = self._updateheight()
67     newroot._height = newroot._updateheight()
68     return newroot

```

## Adding Intervals

When an instance of the interval tree is created, it initially only has a leaf node. This means that new nodes need to be added to represent the endpoints. Adding nodes is a recursive process with many different cases. For this section, assume the interval being added is named `i` and covers  $[start, end]$ .


First of all, if the interval to be added is equal to the interval represented by the current node, then the interval can be added to that node's set. If this condition is met, there's no need to check the cases described below.

After this, the cases can be split up into those that apply at leaf nodes and those that apply at non-leaf nodes.


### Current Node is a Leaf

Note that when converting a leaf node to a non-leaf node, you need to make sure that both children are created.


1. **The current node's interval contains both *start* and *end*:** In this case, the current node is given the boundary value *start*. Its right child is a node with boundary value of *end*. Recursively call the add method on this newly created right child. This ends up adding `i` to the node highlighted below.



2. **The current node's interval contains *start* but *end* is either outside or on the edge:** In this case, the current node is given the boundary value *start*. The right child of the current node becomes a leaf containing `i`.



3. **The current node's interval contains *end* but *start* is either outside or on the edge:** In this case, the current node is given the boundary value *start*. The left child of the currently node becomes a leaf containing `i`.



### Current Node is not a Leaf

Both of these conditions are checked:

1. If the start is less than the boundary, recursively call the add method on the left child. Otherwise, recursively call the add method on the right child.
2. If the end is greater than the boundary, recursively call the add method on the right child. Otherwise, recursively call the add method on the left child.

After adding the nodes, the tree is then rebalanced. Putting all of this together, here is the add method:

```

1  def add(self, start, end, name):
2      # The case that applies to all nodes
3      if start <= self.min and end >= self.max:
4          self.intervals.add(name)
5
6      # Leaf cases
7      elif self._isleaf:
8          # Cases 1 and 2
9          if start > self.min:
10             self.boundary = start
11             # If it's not a leaf I need to make sure it has both children
12             self.left = IntervalNode(min = self.min, max = self.boundary)
13             # Case 1
14             if end < self.max:
15                 self.right = IntervalNode(end,
16                                           min = self.boundary,
17                                           max = self.max)
18                 self.right = self.right.add(start, end, name)
19             # Case 2
20             else:
21                 self.right = IntervalNode(None,
22                                           {name},
23                                           min = self.boundary,
24                                           max = self.max)
25
26             # Case 3
27             elif end < self.max:
28                 self.boundary = end
29                 self.left = IntervalNode(None,
30                                           {name},
31                                           min = self.min,
32                                           max = self.boundary)
33                 self.right = IntervalNode(min = self.boundary, max = self.max)
34
35         # Non-leaf cases

```



```

35     else:
36         # Case 1
37         if start < self.boundary or end <= self.boundary:
38             self.left = self.left.add(start, end, name)
39         # Case 2
40         if start >= self.boundary or end > self.boundary:
41             self.right = self.right.add(start, end, name)
42
43     # Balancing operations
44     if not self._isleaf:
45         self.left._height = self.left._updateheight()
46         self.right._height = self.right._updateheight()
47     self._height = self._updateheight()
48
49     return self.rebalance()


```

## Removal


For removal, you traverse the tree, ending each path when it finds a node containing the interval to be removed or until it reaches a leaf. The paths that need to be traversed are narrowed down based on the endpoints of the given interval.

Removing an interval can potentially leave unnecessary nodes in the tree, though. This implementation has six different cases to remove unnecessary nodes:


1. If a node has two empty non-leaf children, it can be replaced with a leaf that contains the same intervals as the parent.




2. If a node has two identical children, that node can be replaced by either child.




3. If the right child and the left-right grandchild are identical, the current node can be replaced with its left child.




4. The mirror of number three: If the left child and right-left grandchild are identical, the current node can be replaced with its right child.



5. If the current node's left child has no intervals and it's left-right and right-left grandchildren have the same intervals, the tree can be replaced as shown.



6. The mirror of number five: This time, the right child has no intervals.



After checking these conditions, the tree is then rebalanced. Here it is in code, with a few extra utility methods:

```

1  # Case 1
2  @property
3  def _emptychildren(self):
4      return not self._isleaf and \
5             self.left._isleaf and \
6             self.right._isleaf and \
7             self.left.intervals == self.right.intervals
8
9  # Case 2
10 @property
11 def _samechildren(self):
12     return not self._isleaf and \
13            self.left._isleaf and \
14            self.right._isleaf and \
15            self.left.intervals == self.right.intervals
16
17 # Case 3
18 @property
19 def _canreplacewithleft(self):
20     return len(self.intervals) == 0 and \
21            not self._isleaf and \
22            not self.left._isleaf and \
23            self.right._isleaf and \
24            self.right.intervals == self.left.right.intervals
25
26 # Case 4
27 @property
28 def _canreplacewithright(self):
29     return len(self.intervals) == 0 and \
30            not self._isleaf and \
31            not self.right._isleaf and \
32            self.left._isleaf and \
33            self.right.left.intervals == self.left.intervals
34
35 # Case 5
36 @property
37 def _canrestructureright(self):

```

```

38         return not self._isleaf and \
39             not self.left._isleaf and \
40             not self.right._isleaf and \
41             len(self.left.intervals) == 0 and \
42             self.left.right.intervals == self.right.left.intervals and \
43             self.left.right._isleaf and \
44             self.right.left._isleaf
45
46     # Case 6
47     @property
48     def _canrestructureleft(self):
49         return not self._isleaf and \
50             not self.left._isleaf and \
51             not self.right._isleaf and \
52             len(self.right.intervals) == 0 and \
53             self.left.right.intervals == self.right.left.intervals and \
54             self.left.right._isleaf and \
55             self.right.left._isleaf
56
57     def remove(self, interval, start, end):
58         if interval in self.intervals:
59             self.intervals.remove(interval)
60         elif not self._isleaf:
61             if self.boundary >= start:
62                 self.left = self.left.remove(interval, start, end)
63             if self.boundary <= end:
64                 self.right = self.right.remove(interval, start, end)
65
66     # Removing unnecessary nodes
67     if self._emptychildren:
68         self.__init__(None, self.intervals, min = self.min, max = self.max)
69     if self._samechildren:
70         self.__init__(None, self.left.intervals, self.min, self.max)
71     if self._canreplacewithleft:
72         newl = self.left.left
73         newr = self.left.right
74         self.__init__(self.left.boundary, self.left.intervals, \
75             min = self.min, max = self.max)
76         self.left = newl
77         self.right = newr
78     if self._canreplacewithright:
79         newl = self.right.left
80         newr = self.right.right
81         self.__init__(self.right.boundary, self.right.intervals, \
82             min = self.min, max = self.max)
83         self.left = newl
84         self.right = newr
85     if self._canrestructureright:
86         newl = self.left.left
87         newr = self.right
88         self.__init__(self.left.boundary, self.intervals, self.min, self.max)
89         self.left = newl
90         self.right = newr

```

```

91         if self._canrestructureleft:
92             newl = self.left.left
93             newr = self.right
94             self.__init__(self.left.boundary, self.intervals, self.min, self.max)
95             self.left = newl
96             self.right = newr
97
98         # Balancing operations
99         if not self._isleaf:
100             self.left._height = self.left._updateheight()
101             self.right._height = self.right._updateheight()
102             self._height = self._updateheight()
103
104         return self.rebalance()

```

## Testing Points

The process for testing point intersections is simple:

If the given point is in the current node, add the intervals in that node to the return value.

Then, if the point is less than or equal to the boundary, recursively check the left child. If it's greater than or equal, recursively check the right child.

```

1  def testPoint(self, point):
2      ret = set()
3      if self.min <= point <= self.max:
4          ret |= self.intervals
5      if not self._isleaf:
6          if point <= self.boundary:
7              ret |= self.left.testPoint(point)
8          if point >= self.boundary:
9              ret |= self.right.testPoint(point)
10     return ret

```

## Testing Ranges

The process for testing intersection with a range is similar to the process for testing with a point:

```

1  def testRange(self, start, end):
2      ret = set()
3      if self.max >= start and self.min <= end:
4          ret |= self.intervals
5      if not self._isleaf:
6          if start <= self.boundary:
7              ret |= self.left.testRange(start, end)
8          if self.boundary <= end:
9              ret |= self.right.testRange(start, end)
10     return ret

```

## Analysis

The running time of testing a point using this interval tree is  $O(\log n)$ , with  $n$  being the total number of nodes in the tree. At most, the `testPoint` method will traverse paths down the tree if the given point is equal to the root node's boundary.

The running time of testing a range is  $O(n)$ . If the given interval is negative infinity to infinity and the interval to remove is in a leaf node, the function will traverse the whole tree.

Rotation takes  $O(k)$ , where  $k$  is the number of intervals in the nodes that are moved around. This is because those intervals need to be moved between nodes. Because of this, balancing the tree has the same runtime.

Adding an intervals takes  $O(k + \log n)$ . The  $k$  comes from rotations and the  $\log n$  comes from traversing the tree to add to the nodes. At most, the add method will add two new non-leaf nodes.

The running time of removal is  $O(k \log n)$ . The method traverses paths down the tree which have length  $\log n$  and at each node it takes time based on the number of intervals in each node.

Clearing the tree and getting endpoints are constant time operations, as are all other methods in the `IntervalNode`.